



SIGGRAPH THINK
BEYOND
2020 [S2020.SIGGRAPH.ORG](https://s2020.siggraph.org)



JUNE 2020 WEBINAR

Hands-on Workshop: Machine Learning and
Neural Networks

SIGGRAPH NOW



RAJESH SHARMA

SOFTWARE ENGINEER

Walt Disney Animation Studios



Machine Learning

————— Rajesh Sharma —————

Today

- Recap
 - Recurrent Neural Networks
- Reinforcement Learning
- Summary & Resources

Questions?

How are the weights selected for dropout? Is it random?

Greg Klar

The weights of the GAN are being modified in the gradient descent part of the training step, is that correct?

Bobby Bodenheimer

What forces the generator to diversify?

Greg Klar

How to make sure that our discriminator is not "too powerful" at the beginning of the training?

Ivan Puhachov

Why use GAN instead of VAE? Don't the VAE well behaved latent space offer more control over the generation of new samples?

Anonymous Attendee

Can you generate 3D objects with them ?

Ingeborg Tastl

Is an RNN also useful for generating video?

Marijn Eken

Does it matter what algorithm is used to convert text into numbers ?

Ingeborg Tastl

What are RNNs?

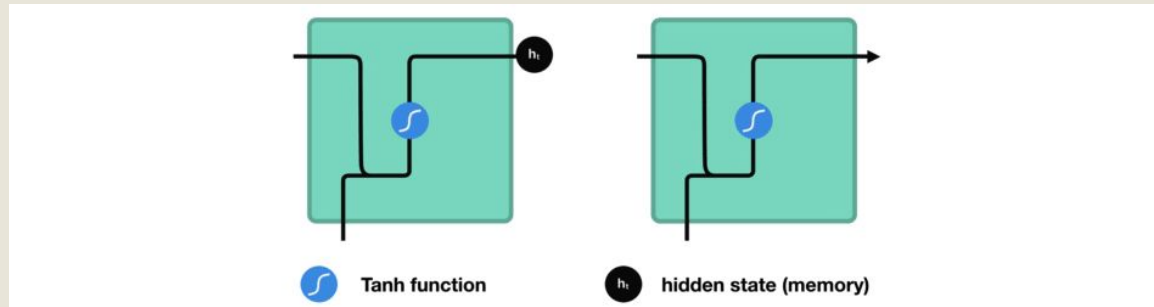
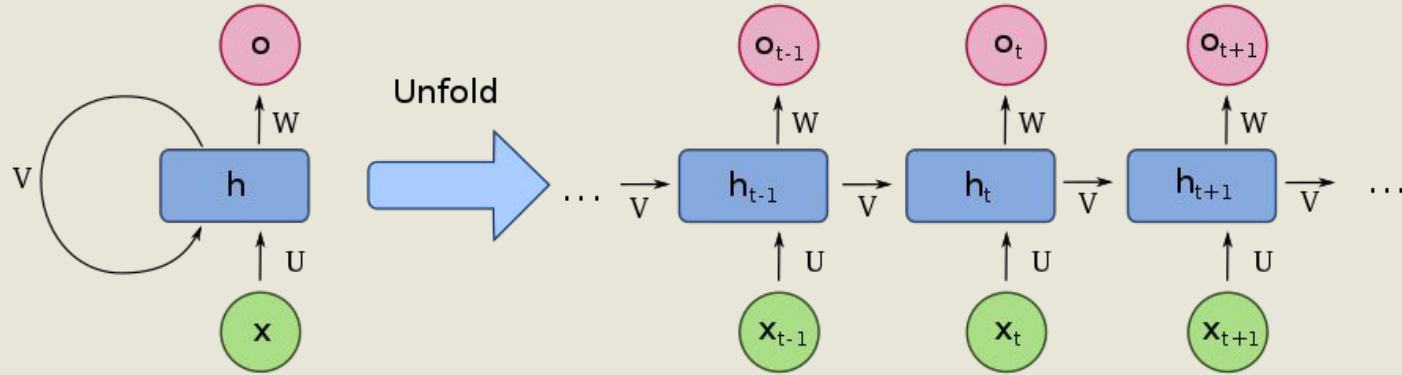
Information from previous timestep is passed forward

Useful for time-dependent data:

- Sequence prediction problems
- Language Translation: Speech, text, music

What are RNNs?



Information from previous timestep is passed forward



Today:

Reinforcement Learning:

Environment, states, actions, rewards



Goal: Maximize Total Reward

Challenges:


Actions & States

Model of environment


Policy for choosing next action

Reward at each step

Value Function



Goal: Maximize Total Reward



N-arm bandit problem




10-arm bandit problem:

- Given 10 choices for action
- Each action has a fixed* reward (unknown)

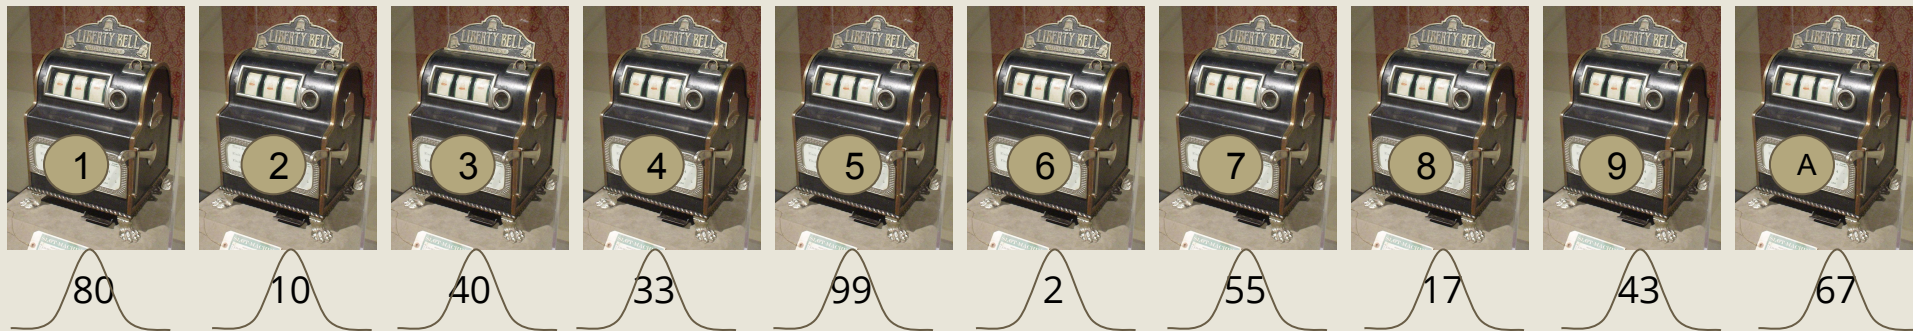
*Reward for each action varies but
each action has a 'fixed' mean reward

Goal: Maximize total reward over a number of actions

10-arm bandit problem:



10-arm bandit problem:



Strategy 1: Naive-Randomly pick an action

Gets an average reward over all actions

Hands on...

Find and open:
tenArmBandit.ipynb

10-arm bandit problem:

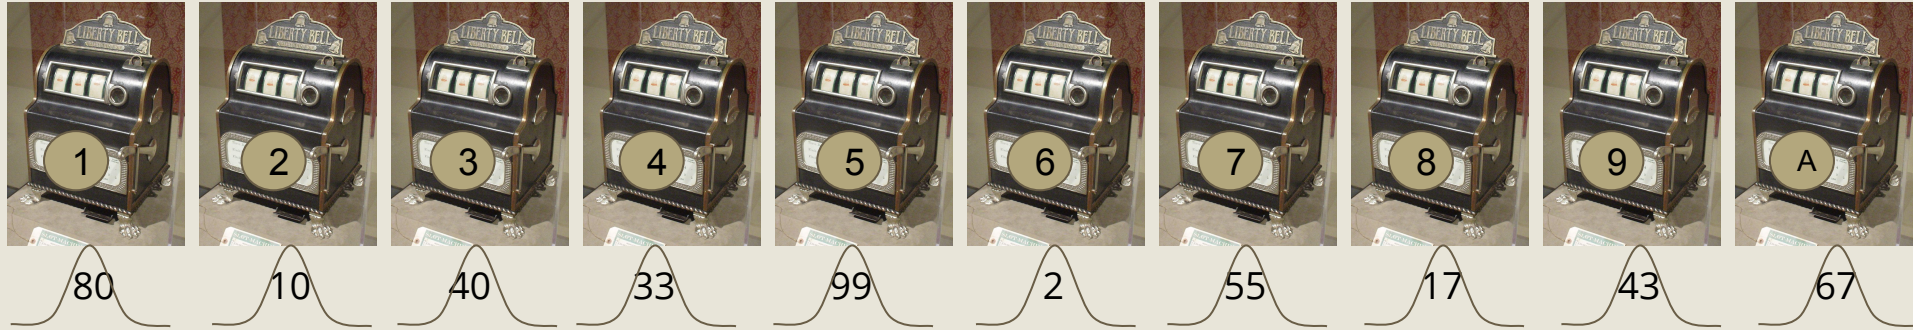
Strategy 2: Greedy

Keep a tally of actions and average reward

Choose action with maximum current avg reward

Q-Action Table

Strategy 2: Greedy



10-arm bandit problem:

Strategy 3: Greedy-epsilon (exploit and explore)

Keep a tally of actions, average reward

Choose action with maximum current avg reward

But,

Once in a while,

choose another action randomly

10-arm bandit problem:

Strategy 4: Greedy-decaying-epsilon

- Keep a tally of actions, average reward
- Choose action with maximum current avg reward

But,

- Once in a while,
 - choose another action randomly
- over time, do this less frequently

Hands on...

Find and open:
tenArmBandit.ipynb

Deep Q Learning (DQN)

- DQN: Deep CNN to represent/learn Q
- Experience Replay: uses a random sample of prior actions instead of the most recent action to proceed
- Example: Space Invaders
- Example: [Walking](#)

Policy Design, reward structure is important



Summary

- Basics: regression, UAT, no free-lunch
- Fully Connected: experiments, final layer
- CNN: building block for image-based training
- RNN, LSTM, GRU: time series, language, text
- Unet, resNet: CNN-like with better detail transfer
- Variational AutoEncoder: Generative:(mean,variance)
- Transfer Learning: mt-cnn, facenet
- GAN: Generative: direct sample
- Reinforcement Learning: env, states, actions, rewards

Summary (Practice)

- Bias vs Variance: Overtraining, Undertraining
- Data: Lots of it, augment, un-biased
- Hyperparameters: layers, nodes, optimizer, L-Rate
- Loss function: logits (log-likelihood), L2, L1
- Training: epochs, batches, tfds, plotting
- Distributions: you are trying to find a sample
- Work like a scientist:
 - Hypothesis, experiment, observe, record, change:
 - Repeat

Local Install (Linux/Mac)

In Colab (after you have imported everything you need)

```
!pip freeze > requirements.txt
```

```
-----  
mkdir ~/myml
```

```
cd ~/myml
```

```
# copy the requirements.txt file from colab to ~/myml
```

```
# save these instructions as README.env in ~/myml
```

```
# create a python3 virtual environment
```

```
python3 -mvenv --system-site-packages mlenv
```

```
# activate the virtual environment
```

```
source mlenv/bin/activate
```

```
# upgrade the installer
```

```
pip install --upgrade pip
```

```
# Install all the software specified in requirements.txt
```

```
pip install -r requirements.txt
```

```
# test it out:
```

```
which python (this should point to the python in the virtual env)
```


Summary (not covered)

- Cloud-based: Training and Deployment
- Local: clusters, machines, environment
- Tensorboard: for logging, visualizations
- Intermediate layer visualization
- Other methods: Random Forests, XGBoost
- More theory

Where to go from here:

Deep Learning: <https://www.deeplearningbook.org>
and some excellent lectures to go along:
https://www.deeplearningbook.org/lecture_slides.html

Reinforcement Learning:
<http://incompleteideas.net/book/bookdraft2017nov5.pdf>

Statistics:
<https://link.springer.com/book/10.1007/978-0-387-21736-9>

A roadmap to reading:
<https://github.com/floodsung/Deep-Learning-Papers-Reading-Roadmap>

More comprehensive list of resources:
<https://www.kdnuggets.com/2020/03/24-best-free-books-understand-machine-learning.html>

Video Tutorials (3Blue1Brown):
https://www.youtube.com/channel/UCYO_jab_esuFRV4b17AJtAw/videos

QUESTIONS?

Submit now!

THANK YOU

WANT TO HEAR MORE FROM SIGGRAPH AND DISNEY?

Check out SIGGRAPH Now on
YouTube



[YouTube.com/user/ACMSIGGRAPH](https://www.youtube.com/user/ACMSIGGRAPH)

Subscribe to the SIGGRAPH Spotlight
podcast

